

Chapter 2: The Basic ER Diagram—A Data Modeling Schema

This chapter begins by describing a data modeling approach and then introduces entity relationship (ER) diagrams. The concept of entities, attributes, relationships, and keys are introduced. The first three steps in an ER design methodology are developed. Step 1 begins by building a one-entity diagram. Step 2 concentrates on using structured English to describe a database. Step 3, the last section in this chapter, discusses mapping the ER diagram to a relational database. These concepts — the diagram, structured English, and mapping — will evolve together as the book progresses. At the end of the chapter we also begin a running case study, which will be continued at the ends of the subsequent chapters.

What Is a Data Modeling Schema?

A data modeling schema is a method that allows us to model or illustrate a database. This device is often in the form of a graphic diagram, but other means of communication are also desirable — non computer-people may or may not understand diagrams and graphics. The ER diagram (ERD) is a graphic tool that facilitates data modeling. The ERD is a subset of "semantic models" in a database. Semantic models refer to models that intend to elicit meaning from data. ERDs are not the only semantic modeling tools, but they are common and popular.

When we begin to discuss the contents of a database, the data model helps to decide which piece of data goes with which other piece of data on a conceptual level. An early concept in databases is to recognize that there are levels of abstraction that we can use in discussing databases. For example, if we were to discuss the filing of "names," we could discuss this:

Abstractly, that is, "we will file names of people we know."

or

Concretely, that is, "we will file first, middle, and last names (20 characters each) of people we know, so that we can retrieve the names in alphabetical order on last name, and we will put this data in a spreadsheet format on package x."

If a person is designing a database, the first step is to abstract and then refine the abstraction. The longer one stays away from the concrete details of logical models (relational, hierarchical, network) and physical realizations (fields [how many characters, the data type, etc.] and files [relative, spreadsheet]), the easier it is to change the model and to decide how the data will eventually be physically realized (stored). When we use the term "field" or "file," we will be referring to physical data as opposed to conceptual data.

Mapping is the process of choosing a logical model and then moving to a physical database file system from a conceptual model (the ER diagram). A physical file loaded with data is necessary to actually get data from a database. Mapping is the bridge between the design concept and physical reality. In this book we concentrate on the relational database model due to its ubiquitousness in contemporary database models.

What Is an Entity Relationship (ER) Diagram?

The ER diagram is a semantic data modeling tool that is used to accomplish the goal of abstractly describing or portraying data. Abstractly described data is called a ***conceptual model***. Our conceptual model will lead us to a "schema." A ***schema*** implies a permanent, fixed description of the structure of the data. Therefore, when we agree that we have captured the correct depiction of reality within our conceptual model, our ER diagram, we can call it a schema.

An ER diagram could also be used to document an existing database by reverse-engineering it; but in introducing the subject, we focus on the idea of using an ER diagram to model a to-be-created database and deal with reverse-engineering later.

Defining the Database — Some Definitions: Entity, Relationship, Attribute

As the name implies, an ER diagram models data as **entities** and **relationships**, and entities have **attributes**. An **entity** is a thing about which we store data, for example, a person, a bank account, a building. In the original presentation, Chen (1976) described an entity as a "thing which can be distinctly identified." So an entity can be a person, place, object, event, or concept about which we wish to store data.

The name for an entity must be one that represents a type or class of thing, not an instance. The name for an entity must be sufficiently generic but, at the same time, the name for an entity cannot be too generic. The name should also be able to accommodate changes "over time." For example, if we were modeling a business and the business made donuts, we might consider creating an entity called DONUT. But how long will it be before this business evolves into making more generic pastry? If it is anticipated that the business will involve pastry of all kinds rather than just donuts, perhaps it would be better to create an entity called PASTRY — it may be more applicable "over time."

Some examples of entities include:

- Examples of a person entity would be EMPLOYEE, VET, or STUDENT.
- Examples of a place entity would be STATE or COUNTRY.
- Examples of an object entity would be BUILDING, AUTO, or PRODUCT.
- Example of an event entity would be SALES, RETURNS, or REGISTRATION.
- Examples of a concept entity would be ACCOUNT or DEPARTMENT.

In older data processing circles, we might have referred to an entity as a record, but the term "record" is too physical and too confining; "record" gives us a mental picture of a physical thing and, in order to work at the conceptual level, we want to avoid device-oriented pictures for the moment. In a database context, it is unusual to store information about one entity, so we think of storing collections of data about entities — such collections are called **entity sets**. Entity sets correspond to the concept of "files," but again, a file usually connotes a physical entity and hence we abstract the concept of the "file" (entity set) as well as the concept of a "record" (entity). As an example, suppose we have a company that has customers. You would imagine that the company had a customer entity set with individual customer entities in it.

An entity may be very broad (e.g., a person), or it may be narrowed by the application for which data is being prepared (like a student or a customer). *Broad* entities, which cover a whole class of objects, are sometimes called generalizations (e.g., person), and *narrower* entities are sometimes called specializations (e.g., student). In later diagrams (in this book) we will revisit generalizations and specializations; but for now, we will concern ourselves with an application level where there are no subgroups (specializations) or supergroups (generalizations) of entities.

When we speak of capturing data about a particular entity, we refer to this as an *instance*. An entity instance is a single occurrence of an entity. For example, if we create an entity called TOOL, and if we choose to record data about a screwdriver, then the screwdriver "record" is an instance of TOOL. Each instance of an entity must be uniquely identifiable so that each instance is separate and distinctly identifiable from all other instances of that

type of entity. In a customer entity set, you might imagine that the company would assign a unique customer number, for example. This unique identifier is called a [key](#).

A [relationship](#) is a link or association between entities. Relationships are usually denoted by verb phrases. We will begin by expanding the notion of an entity (in this chapter and the next), and then we will come back to the notion of a relationship (in [Chapter 4](#)) once we have established the concept of an entity.

An [attribute](#) is a property or characteristic for an entity. For example, an entity, AUTOMOBILE, may have attributes type, color, vehicle_id, etc.

A Beginning Methodology

Database modeling begins with a description of "what is to be stored." Such a description can come from anyone; we will call the describer the "user." For example, Ms. Smith of Acme Parts Company comes to you, asking that you design a database of parts for her company. Ms. Smith is the user. You are the database designer. What Ms. Smith tells you about the parts will be the database description.

As a starting point in dealing with a to-be-created database we will identify a central, "primary" entity — a category about which we will store data. For example, if we wanted to create a database about students and their environment, then one entity would be STUDENT (our characterization of an entity will always be in the singular). Having chosen one first primary entity, STUDENT, we will then search for information to be recorded about our STUDENT. This methodology of selecting one "primary" entity from a data description is our first step in drawing an ER diagram, and hence the beginning of the requirements phase of software engineering for our database.

Once the "primary" entity has been chosen, we then ask ourselves what information we want to record about our entity. In our STUDENT example, we add some details about the STUDENT — any details that will qualify, identify, classify, or express the state of the entity (in this case, the STUDENT entity). These details or contents of entities are called *attributes*.

[1] Some example attributes of STUDENT would be the student's name, student number, major, address, etc. — information about the student.

[1]C. Date (1995) prefers the word "property" to "[attribute](#)" because it is more generic and because "[attribute](#)" is used in other contexts. We will use "[attribute](#)" because we believe it to be more commonly used.

ER Design Methodology

Step 1: Select one primary entity from the database requirements description and show attributes to be recorded for that entity.

"Requirements definition" is the first phase of software engineering where the systems analyst tries to find out what a user wants. In the case of a database, an information-oriented system, the user will want to store data. Now that we have chosen a primary entity and some attributes, our task will be to:

- Draw a diagram of our first-impression entity (our primary entity).
- Translate the diagram into English.
- Present the English (and the diagram) back to the user to see if we have it right and then progress from there.

The third step is called "*feedback*" in software engineering. The process of refining via feedback is a normal process in the requirements/specification phases. The feedback loop is essential in arriving at the reality of what one wants to depict from both the user and analyst viewpoints. First we will learn how to draw the entity and then we will present guidelines for converting our diagram into English.

Checkpoint 2.1

1. Of the following items, determine which could be an entity and state why: automobile, college class, student, name of student, book title, number of dependents.
2. Why are entities not called files or records?
3. What is mapping?
4. What are entity sets?
5. Why do we need Entity-Relationship Diagrams?
6. What are attributes? List attributes of the entities you found in question 1 (above).
7. What is a relationship?

A First "Entity-Only" ER Diagram: An Entity with Attributes

To recap our example, we have chosen an example with a "primary" entity from a student information database — the student. Again note that "a student" is something about which we want to store information (the definition of an entity). In this chapter, we do not concern ourselves with any other entities.

Let us think about some attributes of the entity STUDENT; that is, what are some attributes a student might have? A student has a name, an address, and an educational connection. We will call the educational connection a "school." We have picked three attributes for the entity STUDENT, and we have also chosen a generic label for each: name, address, school.

We begin our first venture into ER diagrams with a "Chen-like" model. Chen (1976) introduced the idea of the ER diagrams. He and others have improved the ER process over the years; and while there is no standard ERD model, the Chen-like model and variants thereof are common. After the "Chen-like" model, we introduce other models. We briefly discuss the "Barker/Oracle-like" model later (in [Chapter 10](#)). Chen-like models have the advantage that one does not need to know the underlying logical model to understand the design. Barker models and some other models require a full understanding of the relational model, and the diagrams are affected by relational concepts.

To begin, in the Chen-like model, we will do as Chen originally did and put the entities in boxes and the show attributes nearby. One way to depict attributes is to put them in circles or ovals appended to the boxes — see [Figure 2.1](#) (top and middle). [Figure 2.1](#) (bottom) is an alternative style of depicting attributes. The alternative attribute style ([Figure 2.1](#), bottom) is not as descriptive, but it is more compact and can be used if Chen-like diagrams become cluttered.

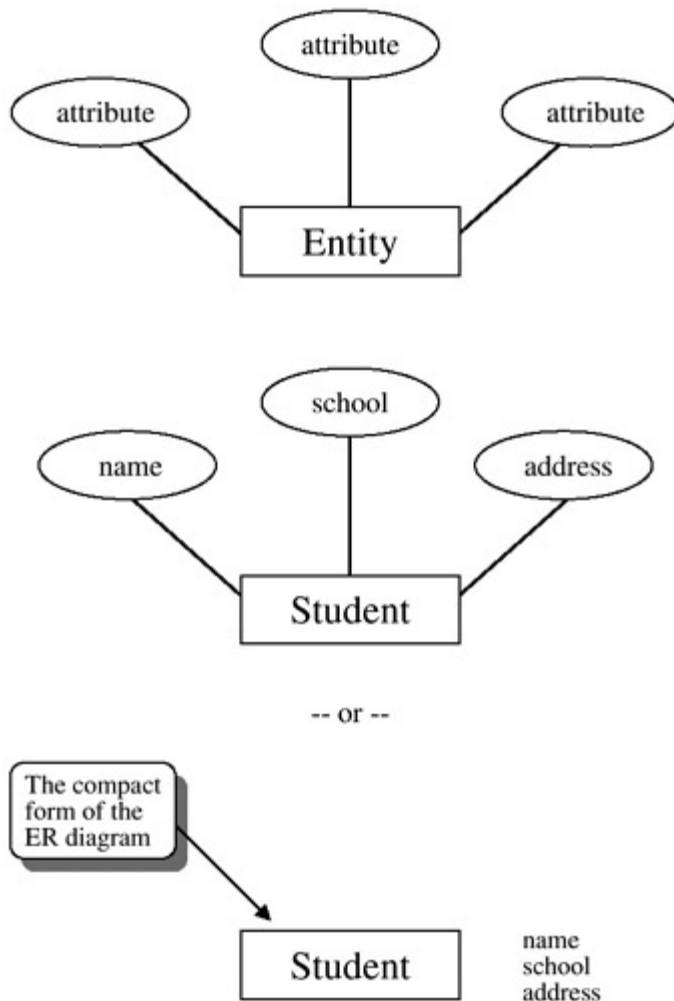


Figure 2.1: An ER Diagram with Three Attributes

There are several ways of depicting attributes. We have illustrated the "attribute in a circle" model (Chen-like model) because it is common and useful. Refer to [Figure 2.2](#) for some alternate models for attributes. There are benefits to alternate forms for depicting attributes. The standard form of the Chen-like model with bubbles and boxes is good for conceptualizing; it is easily changed and very clear as to which attribute goes where. The concise form ([Figure 2.1](#) [bottom] and other variants in [Figure 2.2](#)) is easily created from the standard form and is sometimes more useful for documentation when space is a concern.





Figure 2.2: An ER Diagram with One Entity and Five Attributes, Alternate Models (Batini, Ceri, Navathe)

[Figure 2.1](#) (middle and bottom) shows an ER diagram with one entity, STUDENT, and three attributes: name, address, and school. If more attributes are added to our conceptual model, such as phone and major, they would be appended to the entity (here, STUDENT is the only entity we have), as can be seen in [Figure 2.3](#).




Figure 2.3: An ER Diagram with One Entity and Five Attributes

More about Attributes

Attributes are characteristics of entities that provide descriptive details about the entities. There are several different kinds of attributes: simple or atomic, composite, multi-valued, and derived. The properties of an attribute are its name, description, format, and length, in addition to its atomiticity. Some attributes may be considered unique identifiers for an entity. This section also introduces the idea of a key attribute, a unique identifier for an entity.

The Simple or Atomic Attribute

Simple or atomic attributes cannot be further broken down or subdivided, hence the notion "atomic." One can examine the domain of values^[2] of an attribute to elicit whether an attribute is simple or not. An example of a simple or atomic attribute would be Social Security number, where a person would be expected to have only one, undivided Social Security number.

Other tests of whether an attribute is simple or atomic will depend entirely on the circumstances that the database designer encounters — the desire of the "user" for which the database is being built. For example, we might treat a phone number attribute as simple in a particular database design, but in another scenario we may want to divide the phone number into two distinct parts, that is, the area code and the number. Another example of where the use of the attribute in the database will determine if the attribute is simple or atomic is — a birthdate attribute. If we are setting up a database for a veterinary hospital, it may make sense to break up a birthdate field into month, day, and year, because it will make a difference in treatment if a young animal is five days old versus if it is five months or five years old. Hence, in this case, birthdate would be a composite attribute. For a RACE HORSE database, however, it may not be necessary to break up a birthdate field into month/day/year, because all horses are dated only by the year in which they were born. In this latter case, birthdate, consisting of only the year, would be atomic.

If an attribute is non-atomic, it needs to be depicted as such on the ER diagram. The following sections deal with these more complicated, nonatomic attribute ideas — the composite attribute and the multi-valued attribute.

The Composite Attribute

A composite attribute, sometimes called a group attribute, is an attribute formed by combining or aggregating related attributes. The names chosen for composite attributes should be descriptive and general. The concept of "name" is adequate for a general description, but it may be desirable to be more specific about the parts of this attribute. Most data processing applications divide the name into component parts. Name, then, is called a **composite attribute** or an **aggregate** because it is usually composed of a first name, a last name, and a middle initial — sub-attributes, if you will. The way that composite attributes are shown in ER diagrams in the Chen-like model is illustrated in [Figure 2.4](#). The sub-attributes, such as first name, middle initial, and last name, are called **simple, atomic, or elementary** attributes. The word "aggregate" is used in a different sense in some database query languages and to avoid confusion, but we will not call composite attributes "aggregates;" we will use the word "composite."




Figure 2.4: An ER Diagram with a Composite Attribute — name

Once again, the test of whether or not an attribute will be composite will depend entirely on the circumstances that the database designer encounters — the desire of the "user" for which the database is being built. For example, in one database it may not be important to know exactly which city or state or zip code a person comes from, so an address attribute in that database may not be broken up into its component parts; it may just be called address. Whereas in another database, it may be important to know which city and state a person is from; so in this second database we would have to break up the address attribute into street address, city, state, and zip code, making the address attribute a composite attribute.

The Multi-Valued Attribute

Another type of non-simple attribute that has to be managed is called a **multi-valued attribute**. The multi-valued attribute, as the name implies, may take on more than one value for a given occurrence of an entity. For example, the attribute school could easily be multi-valued if a person attends (or has attended, depending on the context of the database) more than one school. As a counter example, most people go by only one name and hence the grouping, name, is not multi-valued. The multi-valued attribute called school is depicted in [Figure 2.5](#) (Chen-like model) as a double oval, which illustrates the situation where a database will store data about students who may have attended more than one school. Although we have chosen to illustrate school as a multi-valued attribute, we do not mean to imply that this will always be the case in all databases. In fact, the attribute, school, may well be singly valued in some databases. The idea of school may mean the current (or just-previous) school as opposed to all schools attended. If the subjects about whom we are storing data can attend only one school at a time (and that is what we want to depict), then the attribute, school, may well be a single-valued attribute.

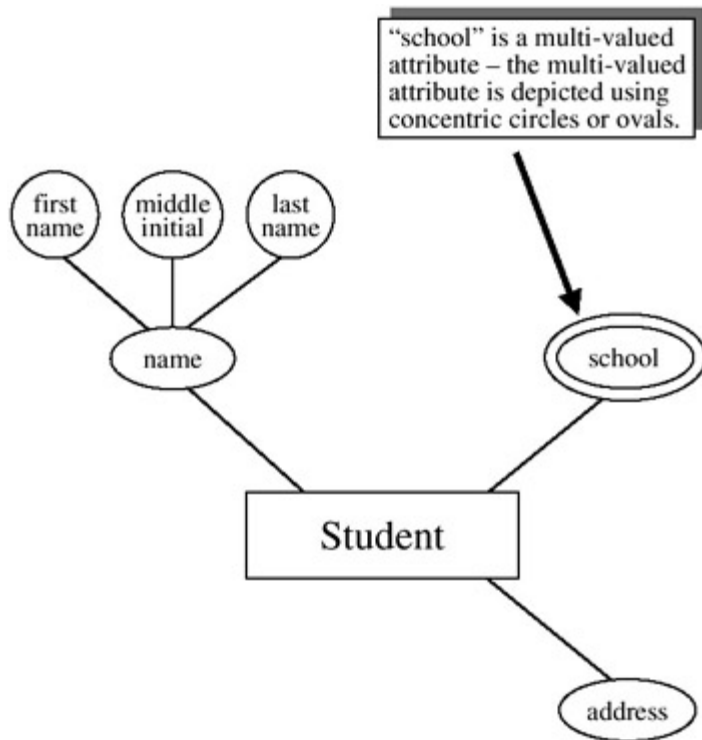


Figure 2.5: An ER Diagram with a Multi-Valued Attribute

Once again, the test of single versus multi-valued will depend entirely on the circumstances that the database designer encounters — the desire of the "user" for which the database is being built. It is recommended that if the sense of the database is that the attribute school means "current school," then the attribute should be called "current school" and illustrated as a single-valued attribute. In our example, we have a multi-valued attribute in [Figure 2.5](#), so the sense of the diagram is that multiple schools can be recorded for each student.

The Derived Attribute

Derived attributes are attributes that the user may envision but may not be recorded per se. These derived attributes can be calculated from other data in the database. An example of a derived attribute would be an age that could be calculated once a student's birthdate is entered. In the Chen-like model, a derived attribute is shown in a dashed oval (as shown in [Figure 2.5A](#)).




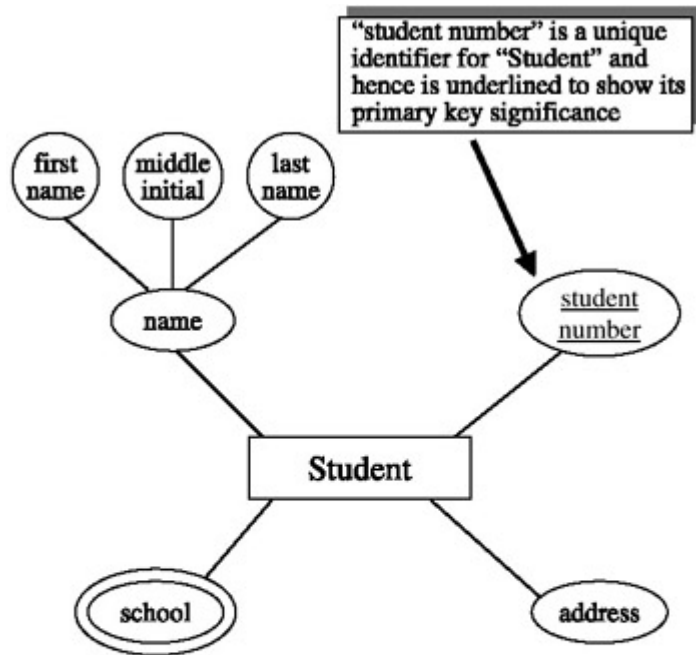
Figure 2.5A: An ER Diagram with a Derived Attribute — age

keys

The sense of a database is to store data for retrieval. An attribute that may be used to find a particular entity occurrence is called a **key**. As we model our database with the ER models, we may find that some attributes naturally seem to be keys. If an attribute can be thought of as a unique identifier for an entity, it is called a **candidate key**. When a candidate key is chosen to be **the** unique identifier, it becomes the **primary key** for the entity.

As an example of keys, suppose we add an attribute called student number to our STUDENT entity example. We might well consider a student number to be a unique identifier for the entity — a candidate key because of uniqueness. Name is often unique, but not necessarily so. Members of the same class often share last names. Address may or may not be a unique identifier and hence is not a likely candidate key. Siblings that take classes together could easily have the same address. The point is that schools often choose to assign a unique student number to each student in order to be able to find student records — the sense of a key is to provide a unique way to find an entity instance (a particular record).

Some schools also choose to record a Social Security number (SSN) as an attribute. An SSN is also unique and hence a candidate key along with student number. If both SSN and student number were recorded, then the designer would have to choose which candidate would be the primary key. In our case, we choose not to record an SSN. The STUDENT entity with the unique identifier student number added as a key, is depicted in [Figure 2.6](#).



or, the short form...




Figure 2.6: An ER Diagram with a Primary Key or Unique Identifier Attribute

In the Chen-like model, attributes that are **unique identifiers** (candidate keys) are usually underlined (as shown in [Figure 2.6](#)). A unique identifier can be an attribute or a combination of attributes. It is not necessary to choose which candidate key will be the primary key at this point, but one could do so. When there is only one candidate key, we will generally speak of it as the primary key, simply because it is obvious that the primary key is a candidate key. In [Figure 2.6](#) we have also depicted the short form of the ER diagram (at the bottom) with composite attributes and multi-valued attributes as well as primary keys. The composite attributes are listed with its component parts, and the multi-valued attributes are enclosed in parentheses in the abbreviated form.

Finally, while on the subject of keys, we will have situations in the ER diagram (in the Chen-like model) where no key is obvious or intended. Entities that have at least one identified key can be called **strong entities**. In Chen's (1976) original article, strong entities were called **regular entities**. Some entities will be discovered which depend on other entities for their being (and hence their identification). Chen called those entities that rely on other entities for their existence, **weak entities**.

We will often be able to recognize these weak entities because they may not have candidate keys, although the actual meaning of a weak entity is "one that depends on another for existence." As Chen did, we will follow the Chen-like notation and call such entities **weak entities** — weak because they will have to depend on some other entity to furnish a unique identifier to

give the entity a reason to be recorded.

Although a weak entity may have a candidate key, it would not be a strong entity. We depict weak entities in the Chen-like ER diagrams with double boxes (see [Figure 2.7](#)). For now, we will concentrate on those entities that have keys, and later we will reconsider situations where no key is obvious.




Figure 2.7: A Strong and a Weak AUTOMOBILE Entity

Checkpoint 2.2

1. Describe the basic types of data representation schemas used in entity–relationship (ER) modeling.
2. What notation is used to diagrammatically show an entity in the Chen-like model?
3. How do we diagrammatically show attributes in the Chen-like model?
4. How do we show composite attributes in the Chen-like model?
5. Draw an entity representation for the entity "building" with the attributes building name, occupancy, and whether or not it has an elevator (yes/no).
6. Embellish the building entity to include the building superintendent's name (first, middle, and last). Does this have to be a composite attribute? Why or why not?
7. Embellish the building entity to include the address of the building,

which will be the primary key.

8. Once again, embellish the building entity to include names (and only the names) of the janitorial staff.
9. Add a multi-valued attribute to the building entity. 10. How many attributes can an entity have?

^[2]The "domain of values" is the set of values that a given attribute may take on. The domain consists of all the possible legal values that are permitted on an attribute. A data type is a broader term used to describe attributes, but "data type" includes the idea of what operations are allowable. Since database people are usually more concerned about storage and retrieval, database "data types" usually just focus on the "domain of values."